



Available online at www.sciencedirect.com

SCIENCE @ DIRECT®

Computers & Industrial Engineering 45 (2003) 545–562

computers &
industrial
engineering

www.elsevier.com/locate/dsw

Hybrid heuristics for planning lot setups and sizes[☆]

Alistair R. Clark*

*Faculty of Computing, Engineering and Mathematical Sciences, University of the West of England,
Coldharbour Lane, Bristol BS16 1QY, UK*

Accepted 8 June 2003

Abstract

The planning of a canning line at a drinks manufacturer is discussed and formulated as a mathematical programming model. Several alternative heuristic solution methods are developed, tested and compared on real data, illustrating the trade-offs between solution quality and computing time. The two most successful methods make hybrid use of local search and integer programming, but in rather different ways. The first method searches for the best proportion by which to factor setup times into unit production times. The second method carries out a local search on the first stage's binary setup variables. In both methods approximate mixed integer programming models are solved at each search iteration. In addition, a local search variant, called diminishing neighbourhood search, is used in order to avoid local optima in a variety of landscapes. Computational tests analyse the quality/time trade-offs between alternative heuristics, enabling an efficient frontier of non-dominated solutions to be identified.

© 2003 Elsevier Ltd. All rights reserved.

Keywords: Production planning; Lot sizing; Setups; Heuristics; Optimisation; Local search

1. Introduction

This paper explores several different heuristic solution approaches for a prototype mixed integer linear programme (MIP) model developed to assist a drinks manufacturer plan the canning of a range of liquid products. The company also bottled these products, but selected the canning line as appropriate for a pilot project, since the line was dedicated to canned products which, in turn, were not produced on any other line. Thus the planning and scheduling of the canning line and its products was carried out independently of the company's bottling lines and associated products.

[☆] This manuscript was processed by Area Editor Maged Dessouky.

* Tel.: +44-117-344-3134; fax: +44-117-344-3155.

E-mail address: alistair.clark@uwe.ac.uk (A.R. Clark).

A concern of the company was to assess the impact of sales promotion commitments and rushed orders on both finished inventory and its ability to fulfil customer orders. The intention was that the model would enable the company to explore the impact of possible promotions and customer order scenarios while using capacity as efficiently as possible. The model sought to optimize the fulfillment of customer orders, taking into account forecast demand, the capacity of the canning line, and the limited availability of liquid mixing tanks, as well as a technological constraint that the tanks should never be used less than half-full. The planning (and scheduling) also had to take into account the impact of product changeovers (setups) on the effective capacity of the canning line.

As the forecasts of individual customer orders and product-specific demand were quite error-prone, there was little point carrying out detailed scheduling for more than a week in advance. However, it was useful to have a provisional production plan for thirteen weeks ahead in order to have advance warning of possible backlogs. Accordingly the proposed system had the following outputs:

1. A production schedule for week 1, detailing which product should be produced on the canning line, their batch sizes and in what sequence.
2. A provisional plan showing the amount produced of each product in each of weeks 2–13 in the first three months. Many products will not be produced every week if scarce capacity is to be used efficiently.

If a large number of products had to be sequenced each week, then a mathematical algorithm would almost certainly identify a more efficient production sequence than a human scheduler. However, the straightforward nature of changeover times meant that the sequencing of products was well handled by the experienced scheduler who simply sequenced together products that used the same liquid. Consequently this paper focuses on the planning model.

Since much demand data is forecast and some production parameters are not precisely known, a solution that is optimal for estimated data may actually add less value than a near-optimal solution. Indeed, it will almost certainly be sub-optimal for the actual data. Rather than seek a perfectly optimal solution, the aim was to find a good solution in practicable time. Thus the objective of this paper is to explore and compare the quality and computing time of several alternative solution approaches. Its purpose is to inform both the academic and practitioner industrial engineering community of the interesting and useful results obtained in this particular case.

The structure of the paper is as follows. Section 2 formulates and explains the planning model and then in Section 3 algorithmic solution methods using only mathematical programming are developed and computationally tested. In Section 4 hybrid heuristics combining local search and mathematical programming are explored and also tested. Finally, Section 5 concludes the paper with a comparative discussion of the test results and provides some pointers for further research. The main conclusion is that, for the situation and data of the drinks manufacturer, hybrid approaches comprising both classical mathematical programming and modern local search methods were the most successful in balancing solution quality with computing time.

2. The planning model

The plant canned $P = 41$ distinct products p filled from $L = 14$ liquids l . The production of both liquids and canned products was planned on a rolling horizon basis for each of $T = 13$ consecutive

weeks t . A basic formulation of the planning problem is:

$$\text{Minimise } \sum_{p,t} (I_{pt}^+ + 100I_{pt}^-) \quad (1)$$

such that

$$I_{p,t-1}^+ - I_{p,t-1}^- + x_{pt} - I_{pt}^+ + I_{pt}^- = d_{pt} \quad \forall p, t \quad (2)$$

$$\sum_p ax_{pt} + s \left(\sum_p y_{pt} - 1 \right) + e \left(\sum_l z_{lt} - 1 \right) \leq B_t \quad \forall t \quad (3)$$

$$y_{pt} \leq z_{lt} \quad \forall l, p \in P(l), t \quad (4)$$

$$x_{pt} \leq M_{pt}y_{pt} \quad \forall p, t \quad (5)$$

The decision variables are:

x_{pt}	quantity of product p produced in week t (≥ 0)
I_{pt}^+	stock of product p at the end of week t (≥ 0)
I_{pt}^-	backlog of product p at the end of week t (≥ 0)
y_{pt}	binary variable taking value 1 if canned product p is produced in week t , and 0 otherwise
z_{lt}	binary variable taking value 1 if liquid l is produced in week t , and 0 otherwise.

The parameter and data inputs are:

d_{pt}	demand for product p at the end of week t
I_{p0}^+	current stock of product p
I_{p0}^-	current backlog of product p
B_t	available time on the line in week t
a	canning-line time required to produce one unit of any product, excluding changeovers
s	canning-line setup time needed to changeover between canned products if no change of liquid is involved
e	extra canning-line setup time needed in the changeover between canned products if a change of liquid is involved
$P(l)$	set of canned products that use liquid l
M_{pt}	upper bound on x_{pt} , calculated as B_t/a

The objective function (1) minimises inventories and backorder penalties, the latter having much more weight than the former, so as to flexibly but strongly discourage planned backorders rather than outrightly prohibiting them. The first constraints (2) balance production, inventories and backlogs with demand. The stock I_{pt}^+ and backlog I_{pt}^- cannot both be positive for a given pair (p, t) , an occurrence that is prevented by both having positive coefficients in the objective function (1). Many products have positive (negative) current stocks which, in a data preparation phase, are discounted from (added to) demand in the first period(s), resulting in zero current stocks and what is known as *effective* demand, calculated as follows:

For $t = 1$ to T do {
 Let $d_{pt} = \max(d_{pt} - I_{p0}, 0)$;
 Let $I_{p0} = \max(I_{p0} - d_{pt}, 0)$;
 If $I_{p0} = 0$ then stop;
 } (end of t loop)
 If $I_{p0} > 0$ then do not plan for product p

where I_{p0} is the current stock (positive or negative) of product p .

Constraints (3) ensure that production and setups take place within the available time. The -1 terms in expression (3) reflect the fact that the first product of the week is either a continuation from the previous week or a new setup carried out on the Sunday in which case it does not consume canning line time. Constraints (4) and (5) ensure that if a product is canned then both it and its liquid need to be set up.

The model has been simplified for the purposes of this paper, taking out the availability of overtime and constraints on the time canned products may stay in inventory before shipping. However, a key feature was the inclusion of liquid production, carried out in tanks of a large fixed size F :

$$\sum_{p \in P(l)} r_{lp} x_{pt} = F n_{lt} \quad \forall l, t \quad (6)$$

where r_{lp} is the amount of liquid l required to make one unit of canned product p and n_{lt} is an integer decision variable for the number of tanks of liquid l to produce and have ready in week t for canning.

For certain slow-moving liquids, it is not always economic to produce a complete tankful, and so a continuous variable $u_{lt} \geq 0$ is included to quantify the underutilisation of a single tank of liquid l in week t :

$$\sum_{p \in P(l)} r_{lp} x_{pt} = F(n_{lt} - u_{lt}) \quad \forall l, t \quad (7)$$

The constraints

$$u_{lt} \leq 0.5 \quad \forall l, t \quad (8)$$

are also necessary since it is not technically viable to produce less than about half a tankful of liquid.

The use of the tanks also needs scheduling each week, but to avoid an overly complex model a capacity limit of C_t tankfuls in week t is simply imposed:

$$\sum_l n_{lt} \leq C_t \quad \forall t \quad (9)$$

In addition, the value of n_{lt} is constrained to be zero if liquid l is not produced in week t or to between 1 and C_t if it is:

$$z_{lt} \leq n_{lt} \leq C_t z_{lt} \quad \forall l, t \quad (10)$$

The complete planning model is given by expressions (1)–(5) and (7)–(10). It is this formulation, denoted CLP (Canned Liquid Planning), that must be solved quickly so as to be suitable for operational use with frequent replanning.

3. Algorithmic solutions with only mathematical programming

Exact optimisation is illusory when production parameter values can vary over time and are not well known, demand forecasts are often revised, and upsets such as machine failure are common, necessitating frequent replanning. A more useful outcome may be a quickly obtained plan of good quality. Bearing this in mind, a number of alternative solution approaches are now explored.

3.1. Default branch-&-bound search on the complete MIP model

The first approach is the ‘lazy’ default—simply let an industrial strength branch-&-bound (BB) MIP solver try to find a good solution within a short amount of time.

BB (Pinedo & Chao, 1999; Wolsey, 1998) is a well-established enumeration method for optimally solving MIPs, based on an upside-down tree search. For problems where the integer variables are only binary, each node in the tree has two branches descending from it, one for the value 0 of a binary variable, the other for the value 1, both leading to a node one level below. At each node, an integer feasible solution is sought, using a heuristic, in order to produce an upper bound on the node’s optimal solution and hopefully update the incumbent solution (the best integer feasible solution found so far at any node). If analysis shows that any solution subsequent to that node is either infeasible or has a lower bound that is higher than the incumbent, then the node is fathomed (i.e. eliminated) because it will not lead to an improved feasible solution. The optimal solution is the incumbent after all branches have been fathomed. If the branch and bound search has to be stopped prematurely, then the incumbent provides a best-so-far feasible solution whose distance from optimality is bounded by the smallest lower bound on any unfathomed node.

For problems with small numbers of integer variables, the BB approach is fast. However, as problem size grows, the number of different integer solution combinations explodes exponentially, possibly causing the search to take an impossibly long time to fathom all branches and converge to a provably optimal solution, as is shown to be the case for the problem of this paper.

The complete MIP model describes a big-bucket lot sizing problem with setup times (Belvaux & Wolsey, 1998), a class which is generally not trivial to solve. For the real 13-week data set used in this paper, the model had 715 binary variables (y and z), 182 non-binary integer variables (n), 1781 continuous variables (x , I^+ , I^- and u), 2171 constraints and 7770 non-zero matrix elements after the presolve had eliminated 82 variables, so it is a sizeable problem, although not huge. It is, however, large enough not to be optimally solvable in a reasonable amount of computer time. The feasible integer solutions shown in Table 1 are the incumbent solution after increasing amounts of BB search time on a Sun Enterprise 450 workstation with a 400 MHz Ultrasparc CPU and 510 MB of RAM. The incumbent solution when the BB search ran out of memory after 13 h was 395,823. Note that the lower bound in Table 1 is stationary after 1 min at 339,528, about 15% below the final incumbent solution. Observe too that the solution does not improve much after the first 30 min of the search.

Clearly the identification of a provably optimal solution plan takes an impracticable amount of computer time and memory, motivating the development of the alternative approaches below.

Table 1
Branch-&-bound results for the complete model

Search time	Best solution (incumbent)	Best lower bound	Branch-&-bound nodes searched
1 min	481,713	339,528	2566
2 min	481,713	339,528	5333
5 min	481,713	339,528	12,993
10 min	481,713	339,528	26,429
15 min	481,713	339,528	38,661
20 min	430,532	339,528	51,600
30 min	414,358	339,528	76,024
1 h	414,358	339,528	136,471
1.5 h	400,408	339,528	201,930
5 h	400,408	339,528	616,063
10 h	395,823	339,528	1,198,559
Out of memory	395,823	339,528	1,527,170

3.2. Accelerated branch-&-bound search on the complete MIP model

An alternative to searching for a provably optimal solution in a BB search is to ruthlessly prune the tree by fathoming a node if its lower bound is greater than the incumbent solution value minus a certain proportion ϕ of the node's upper bound. The BB default corresponds to $\phi = 0$. A positive value of ϕ will generally accelerate the search, but may miss the optimal solution by pruning a branch leading to it. Table 2 shows the results of computational tests with a range of values of ϕ and a time limit of 1 h imposed on the BB search. Note that the good results obtained for values of ϕ between 0.10 and 0.21, and the excellent solution of 367,089 (all stock, no backlogs) found in 720 s (12 min) when $\phi = 0.21$. However, the value of ϕ only needs to be increased slightly to 0.22 for the solution to deteriorate to 466,917 after an hour's BB search.

Table 2
Accelerated branch-&-bound results for the complete model

Value of ϕ	BB solution	Time
0.00	414,358	1 h
0.05	385,002	1 h
0.10	384,534	300 s
0.15	392,153	460 s
0.20	373,452	370 s
0.21	367,089	720 s
0.22	466,917	1 h
0.25	466,917	21 s
0.30	466,917	17 s
0.35	466,917	17 s
0.40	466,917	17 s

This approach provides a useful benchmark as it identifies the best solution (367,089) found in the whole of this paper, but it is too unpredictable in terms of solution quality and computing time to be practical for operational use.

In an attempt to obtain an even better benchmark, the BB search with $\phi = 0$ was repeated with the 367,089 solution specified at the start as the incumbent, but no better solution was identified before the search ran out of memory after 14 h.

3.3. Backward-then-forward pass

The third approach first optimizes production one period at a time in a backward pass through periods $T, T-1, \dots, 1$ in order to identify the target stock levels S_{pt} that would be necessary to fulfill demand after period t :

Backward pass:

Take out the backorder variables I_{pt}^- from formulation CLP.

Let $S_{pT} = 0 \forall p$.

For $t = T$ down to 1 do {

Fix $I_{pt}^+ = S_{pt}$ as data $\forall p$, and free $I_{p,t-1}^+$ as variables $\forall p$.

Solve formulation CLP just for the $I_{p,t-1}^+$ and all period t variables.

Let $S_{p,t-1} = I_{p,t-1}^+ \forall p$.

} (end of t loop).

The target stock levels S_{p0} for period zero provide indicative minimum levels for current stocks I_{p0}^+ that are needed to satisfy demand over the T -week planning horizon without backlogs. If effective demand is used (so that current stocks and backorders are zeroed), then an indicator of insufficient capacity to meet overall demand is the value $\sum_p S_{p0}$.

After the backward pass above, a forward pass through periods $1, 2, \dots, T$ then optimizes production one period at a time in order to build up the target stock levels S_{pt} .

Forward pass:

For $t = 1$ to T do {

Inflate demand by the target stocks levels: let $d_{pt} = d_{pt} + S_{pt} \forall p$.

Solve formulation CLP for just the period t variables.

Restore the original demand: let $d_{pt} = d_{pt} - S_{pt} \forall p$.

Recalculate and then fix stocks $I_{p,t-1}^+$ and backorders $I_{p,t-1}^-$ as data $\forall p$.

} (end of t loop).

The above backward–forward approach is surprisingly effective giving a solution of 504,305 (all stock with zero backlogs) in just 2 s. This fast and practicable computing time includes the solving of $2T = 26$ MIPs, most with 55 binary variables each. If the target stock levels S_{pt} identified in the backward pass are ignored (i.e. set to zero) then, not surprisingly, a far worse solution of 21,948,200 is obtained (with 73,270 stocks and 218,749 backlogs).

3.4. Forward pass with linear setup approximations

A fourth approach involves the elimination of the binary y and z variables representing the setups for week 2 onwards, compensating by increasing the values of the unit production time a in those weeks.

The resulting model, denoted CLP*, had just 55 binary variables, 14 integer variables, 1781 continuous variables, 851 constraints, 4302 non-zero matrix elements and was solved in under a second, i.e. very quickly. After solving for the setups of week 1, the model was reformulated so that the former week 2 was now week 1, and number of weeks in the model reduced by one from T to $T - 1$. The model was solved afresh, and so on, with $T - 2$ further models being solved, until the setups for periods $1, \dots, T$ had all been decided. The whole series of models was solved in between 9 and 12 s, but the quality of the resulting solution depended on how much the value of the unit production time a was increased, as the tests below show.

In model CLP*, constraints (3)–(5) and (7)–(10) all applied only to week 1. For weeks 2 to T , the integer tank n variables were made continuous, and so the tank under-utilisation u variables were eliminated. Thus the constraints for week 2 onwards were:

$$I_{p,t-1}^+ - I_{p,t-1}^- + x_{pt} - I_{pt}^+ + I_{pt}^- = d_{pt} \text{ (unchanged) } \forall p, t \geq 2 \quad (11)$$

$$\sum_p a^* x_{pt} \leq B_t \quad \forall t \geq 2 \quad (12)$$

$$\sum_{p \in P(l)} r_{lp} x_{pt} = F n_{lt} \quad \forall l, t \geq 2 \quad (13)$$

$$\sum_l n_{lt} \leq C_t \quad \forall t \geq 2 \quad (14)$$

where a^* is the unit production time with setups factored in.

The resulting provisional lot-sizes for weeks 2 to T are never actually implemented, but merely used to approximately indicate future production which is thus taken account of in deciding week 1's lot setups and sizes. The actual lot setups and sizes in each of weeks 2 to T are all eventually determined via the week 1 constraints (3)–(5) and (7)–(10) of model CLP, hence ensuring capacity feasibility of production in all of weeks 1 to T .

By how much should a be increased to a^* ? From constraint (3), the value of a^* has to satisfy

$$\sum_p a^* x_{pt} = \sum_p a x_{pt} + s \left(\sum_p y_{pt} - 1 \right) + e \left(\sum_l z_{lt} - 1 \right) \quad \forall t \quad (15)$$

Now, assuming that

1. the total production $\sum_t x_{pt}$ of a product p equals its total effective demand $\sum_t d_{pt}$, which is a reasonable assumption unless overall demand far outstrips capacity,
2. the canning line is using all available capacity time on setups and production, so that B_t is equal to the right hand side of constraint (3) and hence, by Eq. (15), to $\sum_p a^* x_{pt}$

then $\sum_{p,t} a^* d_{pt} = \sum_t B_t$ and so the following expression for the inflated unit production time a^* can be derived

$$a^* = \frac{\sum_t B_t}{\sum_{p,t} d_{pt}} \quad \forall p \quad (16)$$

resulting, for the data in hand, in more than a 54.5% increase from a to a^* .

Using $a^* = 1.545a$ does indeed give a reasonable result of 628,910 (in 9 s), but it is still inferior to that of 504,305 (with zero backlogs) in just 2 s of computing time resulting from the backward-then-forward pass of Section 3.3. Are there values of a^* that produce a better solution? In Fig. 1 the top graph (a) shows how the solution varies for values of a^* over intervals of $10^{-3}a$ between $0.5a$ and $2.5a$. It can

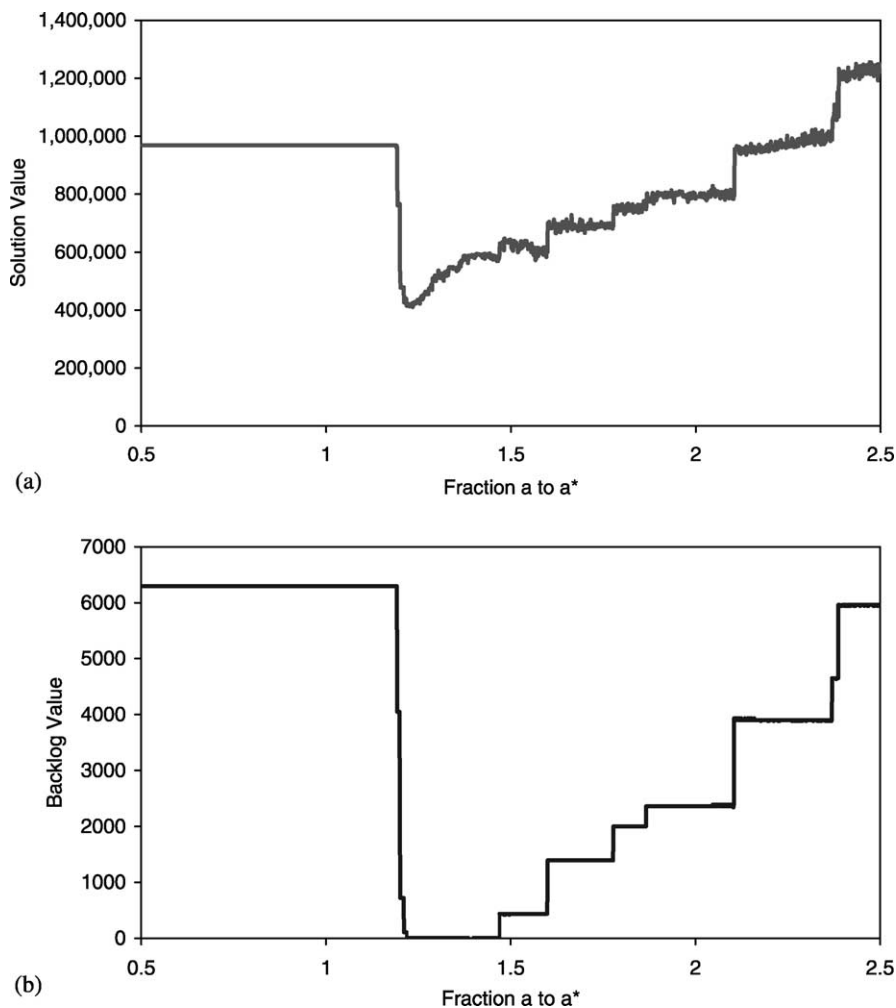


Fig. 1. Solution (a) and backlog (b) values as a function of a^* .

be seen that the minimum solution is around 400,000 and lies somewhere near $1.2a$. The graph also illustrates very clearly that not factoring in setup times, i.e. using $a^* = 1.0a$, results in a solution of about 970,000, well over 100% above the minimum. Note how the solution suddenly falls sharply from 968,823 at $1.189a$ to 477,860 at $1.201a$. Graph (b) in Fig. 1 shows the huge reduction in backlogs, almost to zero, that is achieved by the value of a^* being sufficiently large to force peaks of demand to be produced earlier in time. Note also the subsequent and gradual rise of the solution and backlog values from above $1.5a$ as a^* becomes too large, causing premature preproduction of peaks of demand. It took 5 h and 40 min to calculate all 2000 plotted solutions in Fig. 1, hardly an efficient way of finding a near-optimal solution.

A quicker and more effective method would be to perform an intelligent search for the optimal value of a^* . This is easier said than done, as can be seen from the top graph (a) of Fig. 2 which shows how

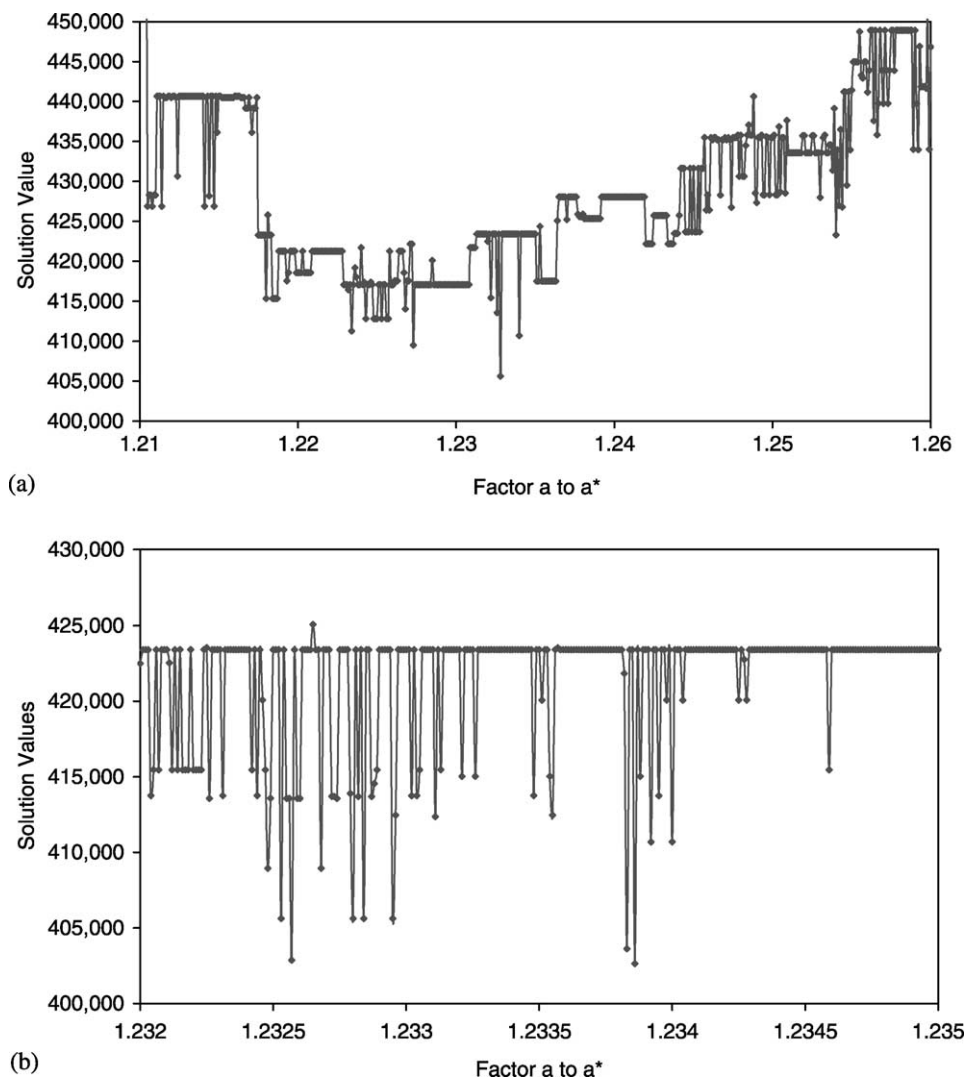


Fig. 2. Solution value as a function of a^* .

the solution varies in a spiky manner for values of a^* over intervals of $10^{-4}a$ between $1.2100a$ and $1.2600a$. The bottom graph (b) reinforces the impression of spikiness of solution values over intervals of $10^{-5}a$ between $1.23200a$ and $1.23500a$. In other words, the solution is sensitive to small changes in the value of a^* which are altering the combinatorial schedule of setups which in turn impact on the availability of canning line time and hence on the stocks and, to a lesser extent in this range, the backlogs. For example, the three values $1.23256a$, $1.23257a$, and $1.23258a$ have, respectively, the solutions 413,571, 40,287 and 423,401, all three of which differ between each other with respect to several of the tank variables z_{lt} (although all three have just 17 units of backlogs). As can be seen on graph (b) of Fig. 2, the third solution value 423,401 occurs frequently over the range $1.232a$ to 1.235 (and beyond). Inspection of these solutions reveals that they differ also slightly with respect to the tank variables z_{lt} , but again all have just 17 units of backlogs as, in fact, do all solutions from 1.21875 to 1.26 .

This spikiness has implications for the way we try to automatically identify the best value of a^* for a given problem instance, as will be discussed in Section 4.

4. Hybrid heuristics with local search and mathematical programming

The final two approaches both use neighbourhood or *local search* (Aarts & Lenstra, 1997; Pirlot, 1996; Sait & Youssef, 1999), a metaheuristic iterative method for finding an optimal or good solution to general combinatorial problems. In the first iteration, a starting solution is identified and adopted as the current solution. At each iteration, the current starting solution is adjusted to produce one or more so-called neighbouring solutions. In the most simple form of local search, known as *random hill climbing* or *ascent algorithm*, if a randomly selected neighbouring solution has an objective function that is an improvement over the current solution, then that neighbour is adopted as the new current solution. If not, then further randomly selected neighbouring solutions are evaluated until an improved one is found. If it is not possible to find a better solution within a certain amount of time or effort, then the local search is stopped and the current solution hailed as the best one found.

The simple form of local search can be likened, in the case of a maximisation problem, to a person trying to walk up a hill blindfolded by feeling around with your feet for a direction with an upward slope and then walking one or more steps in that direction. At the new position, the person feels around again for an upwards slope. In this manner, she or he will eventually reach the top of the local high point, but probably will not reach the top of the highest point in the country, let alone Mount Everest. Thus the simple ascent form of local search will get near a local optimum, but not necessarily near the global optimum, depending on the complexity of the problem, how neighbouring solutions are defined, and the computing effort one is prepared to spend on the search.

Various strategies have been proposed to encourage the local search to not get stuck at a local optimum, but rather to converge to the global optimum, such as simulated annealing (Aarts, Korst, & van Laarhoven, 1997; Dowsland, 1993; Sait & Youssef, 1999), which permits moving to inferior neighbouring solutions with decreasing probability, the related strategy of threshold accepting (Dueck & Scheuer, 1990; Glass & Potts, 1996; Lin, Haley, & Sparks, 1995; Meyr, 2000) where inferior moves are accepted deterministically if they worsen by no more than a decreasing threshold, and tabu search (Glover & Laguna, 1997; Hertz, Taillard, & de Werra, 1997; Sait & Youssef, 1999) where recent moves are banned through a tabu list of temporarily prohibited moves. A good tutorial and comparative discussion of both these and other local search methods can be found in Pirlot (1996).

The local search approach proposed below starts with a very wide range of neighbours in order to generate a diverse range of solutions and avoid premature entrapment at a local optimum. The size of the neighbourhood gradually diminishes as the search progresses in order to home in on a good solution. Moves to an inferior solution are not permitted.

The challenge here is to identify the size of an effective neighbourhood structure at different stages of the search. At one extreme, a small neighbourhood structure could mean a very slow convergence rate, coupled with the risk of getting trapped in a local optimum. At the other extreme, a very wide ranging structure would lean in the direction of pure random sampling among all possible solutions. Random sampling has the advantage that it would avoid entrapment in a local optimum, but is almost certainly inefficient at identifying good solutions by nature of its unguided randomness. A possible compromise is to start a local search with the largest possible neighbourhood to avoid bad or mediocre local optima, and then narrow the neighbourhood as the search progresses in order to home in to a good solution, even if it is still a non-global optimum. The approach is given the name *diminishing-neighbourhood (DN) search*. Other authors (Ahuja, Ergun, Orlin, & Punnen, 1999; Hansen & Mladenovic, 2001; Mühlenbein & Zimmerman, 2000; Reeves, 1994) have advocated the use of large or varying-sized neighbourhoods to avoid local optima, but not in the continuously diminishing manner proposed here.

4.1. Local search on the forward pass with linear setup approximations

The spikiness manifested in Figs. 1 and 2 suggests that some neighbourhood searches on a^* might get stuck in a local optimum far from the global optimum around $1.2a$. How then does diminishing neighbourhood search perform? Computational tests were performed using the following search approach:

1. Start with a reasonable estimate of a^* , such as the value of 1.545 from expression (16). Make this the current solution.
2. Set the neighbourhood radius N to an initial value IN .
3. Iterate \mathcal{N} times. At each iteration:
 - (a) Generate a neighbouring solution by randomly selecting a real number r in the interval $[-N, +N]$ and setting neighbouring $a^* = \text{current } a^* + r$.
 - (b) Solve model CLP* using neighbouring a^* .
 - (c) Move to the neighbouring a^* solution if it is better than the current one.
 - (d) Reduce N by multiplying by a constant factor f .
4. Stop and output the current solution as the best found by the search.

Starting with a good estimated solution (here 1.545), the starting neighbourhood radius IN should be set sufficiently large to be able to encompass the likely optimal solution, although this is not strictly necessary and will sometimes be difficult to judge in practice.

In general, a parameter determining the effectiveness of diminishing neighbourhood search will be the rate at which the neighbourhood size diminishes. If this rate is too small for the number of iterations carried out, then too much time will be spent on random-like search (diversification) and not converge sufficiently (intensification) to a good solution. If the rate is too large, then too much intensification will occur in relation to diversification and the search will prematurely home in to a local optimum which

may be much worse than the global optimum. In this application, N is reduced by being multiplied by a constant factor $f < 1$ at every iteration.

Computational tests were carried out using $\mathcal{N} = 100$ iterations, each search taking about 17 min. Using principally an initial neighbourhood radius $IN = 0.5$, alternative reduction factors $f = (FN/IN)^{1/\mathcal{N}}$ were each tested 50 times, corresponding to a number of finishing radiuses FN , as well as a random search over $1.545 \pm IN$ for benchmarking purposes. Note that $IN = FN$ implies a pure descent algorithm.

The results are shown in Table 3. The parameters IN and FN were initially estimated as 0.5 and 0.0005, respectively, and indeed these performed best, giving a solution of 408,965. Using $\mathcal{N} = 1000$ iterations with the same initial and final radiuses improved slightly on this, providing a solution of 404,549. Several other initial radiuses with $FN = 0.0005$ were also tested over 100 iterations, but with inferior results. The pure descent algorithm and random search performed poorly, indicating the power of DN search to cope with the spikiness of the a^* landscape if appropriate initial and final radiuses are used. It is interesting that an initial radius of $IN = 1.0$ produced a solution only slightly poorer than that for $IN = 0.5$, possibly suggesting that it is safer to err on the side of a larger initial neighbourhood than a smaller one.

4.2. Local search on the binary tank variables

DN and the descent algorithm are now applied to the binary tank variables z_{lt} of model CLP. A neighbouring solution is generated by carrying out a large number of random changes to the 0/1 values of the z variable of the current solution. As the search progresses, the number of such changes needed to generate a neighbour is slowly reduced. At each iteration of the search, a reduced-size MIP is solved, not optimally, but only approximately in order to rapidly obtain a measure of the quality of the solution being evaluated. This enables the search to proceed and identify good solutions relatively quickly. Specifically, the procedure is:

Table 3
Diminishing neighbourhood and descent algorithm on the forward pass with linear setup approximations

No. of iterations	Starting radius N	Finishing radius N	No. of tests	Mean solution	Best solution	a^* of best solution	Mean time (min)
100	1.0	0.0005	50	410,326	402,777	1.233527	16.7
100	0.5	Random S	50	418,833	410,687	1.234001	17.4
100	0.5	0.5	50	419,812	402,512	1.225303	16.5
100	0.5	0.05	50	415,083	402,628	1.233365	16.6
100	0.5	0.005	50	411,458	404,577	1.231862	16.5
100	0.5	0.001	50	410,296	402,857	1.232674	16.5
100	0.5	0.0005	50	408,965	399,112	1.233123	16.5
100	0.5	0.00005	50	424,705	402,628	1.233163	16.8
100	0.2	0.2	50	415,946	408,753	1.222975	16.4
100	0.2	0.0005	50	434,083	402,628	1.233866	16.8
100	0.1	0.0005	50	449,073	405,624	1.232552	17.1
1000	0.5	0.0005	20	404,549	400,390	1.228465	168.6

1. Start with the incumbent solution $\{z_{lt} | \forall l, t\}$ after 1 min of BB search.
2. Set $N = LT = 182$, the number of z_{lt} variables.
3. Iterate \mathcal{N} times. At each iteration:
 - (a) Generate a neighbouring solution by randomly selecting N distinct pairs (l, t) and then randomly deciding if each pair should be toggled by letting $z_{lt} = 1 - z_{lt}$ with probability 0.5. Ensure that at least one randomly chosen pair is toggled. Fix the values of the z_{lt} variables.
 - (b) Solve the model for the integer n_{lt} , binary y_{pt} , and continuous u_{lt} , x_{pt} , I_{pt}^+ and I_{pt}^- variables, using BB search, stopping at the first integer solution.
 - (c) Move to the neighbouring solution if it is better than the current one.
 - (d) Reduce N (as discussed below).
4. Stop and output the current solution as the best found by the search.
5. Fix the values of the z_{lt} variables in the best solution in step (4) and carry out a BB search over the remaining variables for 5 min.

As there are LT z_{lt} variables, all solutions are reachable and equally probable in step 3a when the search is started with an initial value of $N = LT$ (as in step 2). Generally, for any value of N , all solutions that differ by N or fewer 0/1 values are reachable at step (3a) with equal probability.

For the canning line data of this paper, it takes between 0.15 and 0.7 s in step 3b to find the first integer feasible solution. This is a more viable time at each iteration of the search than spending even a minute looking for a better BB solution. However, it does mean that the solution value is generally an approximation to the optimal value at that iteration, inevitably blunting the search, but still effective in that the time saved will permit a vastly greater number of iterations to be carried out.

A parameter determining the effectiveness of the search will be the rate at which the neighbourhood diminishes, i.e. the rate of reduction of N in step 3d. A real-valued parameter $\text{Real}N$, initially set to N , is reduced by being multiplied by a constant factor $f < 1$ at every iteration and then used to let $N = \text{ceiling}(\text{Real}N)$, the integer at or immediately above $\text{Real}N$. If the value of f is too small for the number of iterations carried out, then the search will spend too much time on random-like search (diversification) and not converge sufficiently (intensification) to a good solution. If f is too large, then too much intensification will occur in relation to diversification and the search will prematurely home in to a local optimum which may be much worse than the global optimum. Rather than specify the values of f in the computational tests below, f was calculated as

$$f = (LT)^{-1/q\mathcal{N}} \quad (17)$$

where \mathcal{N} is the number of iterations in the local search and q the percentage of the search after which the smallest neighbourhood of size $N = 1$ was reached in step 3, as illustrated in Fig. 3.

Computational experiments were carried out to compare the method with the other approaches. The number \mathcal{N} of iterations per search was $\mathcal{N} = 1000$ and 10,000. The parameter varied was the percentage q of the search after which the curve floor of $N = 1$ is reached, using test values $q = 0, 25, 50, 75$, and 100%, as shown in parts (a)–(e) of Fig. 3. In addition, for comparison, the tests include the curve (with $f = (0.5)^{1/\mathcal{N}}$) for which $N = LT/2$ at the end of the search as in part (f) of Fig. 3.

One thousand iterations of local search took approximately 10 min while 10,000 iterations took about 10 times as long. At the end of the local search, a further 5 min was spent on a final BB search in a generally successful attempt to improve the solution. Twenty-five 1000-iteration searches were carried

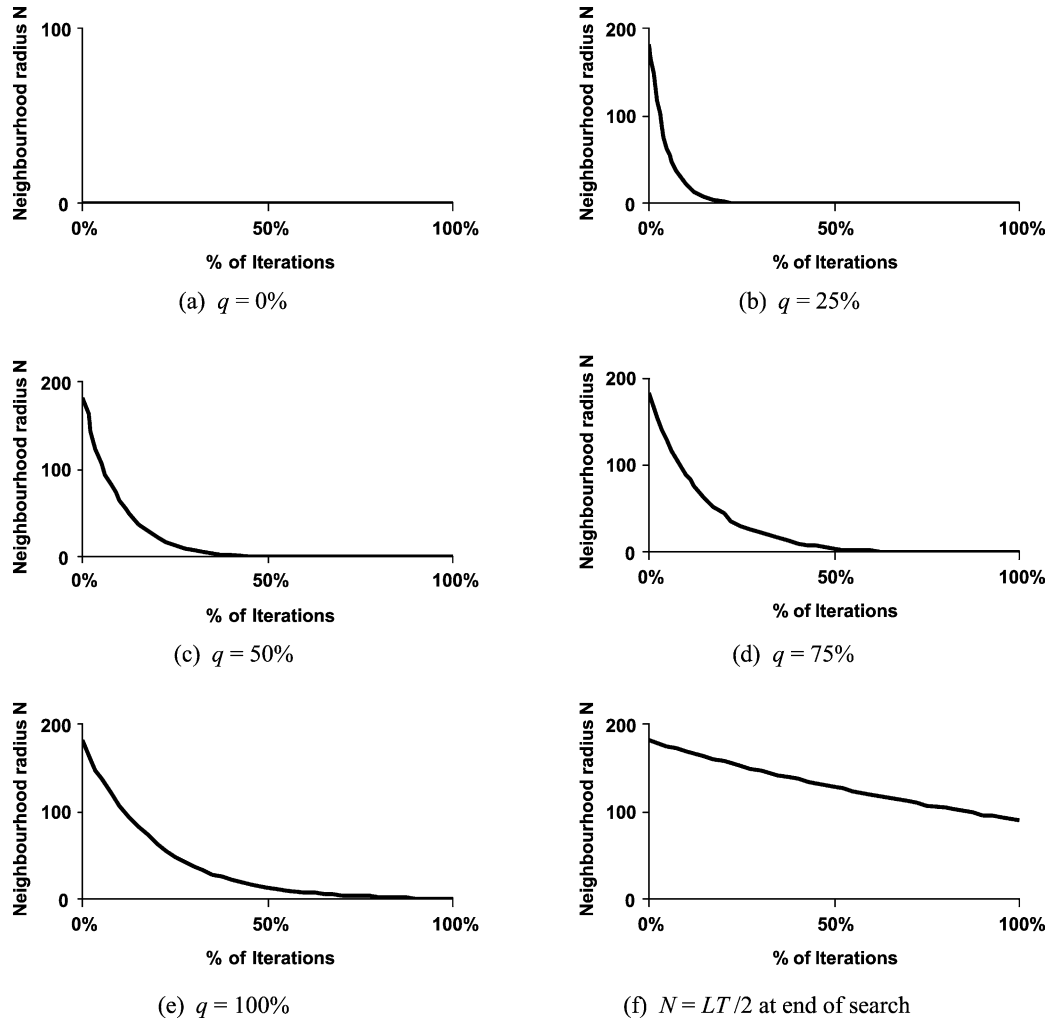


Fig. 3. Reduction rate of the neighbourhood radius during the search.

out for each of the six curves, but due to the larger computing times, only ten 10,000-iteration searches were made for the six curves. The results are shown in Table 4.

For both 1000 and 10,000 iterations, the diminishing neighbourhood method gives its best mean solution values of q between 25 and 100% and appears to be robust within this range. The improvement from 1000 to 10,000 iterations is clear. The curve where $N = LT/2$ at the end of the search results in absolutely no improvement during the search—the solution of 427,088 is due to the final 5 min of BB search improving on the initial solution of 481,713, after fixing the z_{lt} variables. When 1000 iterations are used, the curve $q = 0\%$ (i.e. a simple local search with no diversification at all, as the smallest neighbourhood of $N = 1$ is used throughout) results in a mean solution that is a little better than those for curves $q = 25–100\%$.

However, when 10,000 iterations are used, the mean solution for $q = 0\%$ only slightly improves to around 411,271, resulting in a worse mean solution than for $q = 25–100\%$. This suggests that

Table 4
Diminishing neighbourhood search results

Curve: rate of reduction of N	One thousand iterations 25 runs (16 min each)		Ten thousand iterations 10 runs (110 min each)	
	Mean solution	Best solution	Mean solution	Best solution
Descent: $q = 0\%$	411,855	383,592	411,271	387,743
DNS: $q = 25\%$	414,107	384,230	401,738	376,952
DNS: $q = 50\%$	414,739	388,101	394,006	379,325
DNS: $q = 75\%$	413,199	393,026	402,711	380,977
DNS: $q = 100\%$	416,259	376,287	399,365	388,137
DNS: $N \searrow LT/2$	427,088	427,088	427,088	427,088

the descent algorithm tends to get trapped near a local optimum at between 1000 and 10,000 iterations, as evidenced by the fact that the best solution in a 10,000-iteration search was reached after just 2811 iterations on average. Even the $q = 25\%$ curve, with its limited diversification at the start, enables the diminishing neighbourhood search to avoid mediocre local optima and progress by iteration 3242 on average to a better mean solution of 401,738 (whose constituent solutions are almost certainly local optima, given that the remaining average of 6758 iterations did not identify a better solution from the $LT = 182$ neighbours within the $N = 1$ neighbourhood).

For curves $q = 25\text{--}100\%$, the final 5 min of BB search in step 5 make a mean difference of about $2\frac{1}{2}\%$ for 1000 iterations (improving the mean solution from about 425,000 to around 414,500), and of about $1\frac{1}{2}\%$ for 10,000 iterations (improving the mean solution from 405,300 to 399,500). This suggests that the final 5-min BB search is worthwhile, but not crucial. The gains in the final BB search indicate that some of the improvement potential lost in step 3b by terminating the BB search at the first integer solution is recoverable in step 5. It would be totally recoverable if in step 3b the rank ordering of 5-min solutions and of first integer solutions were the same, thus following identical search trajectories. From the perspective of optimality, the implicit assumption in step 5 that the rankings of the first integer solutions and optimal solutions are positively correlated is not unreasonable and, while blunting the search, certainly makes it viable in reasonable computing time.

5. Discussion and conclusions

Table 5 summarises the mean quality and computing time of the solutions obtainable in the operational use of the heuristics. The results of Table 2 are thus excluded from consideration because they offer no guidance as to what is a good (let alone best) value of ϕ that can be expected a priori to give good solutions. The value of ϕ has only to shift slightly from 0.21 to 0.22 for the solution to change from the best in this paper (367,089 in 720 s) to a mediocre one (466,917 in 1 h). In Tables 3 and 4, DN search showed itself to be more robust with respect to its parameters.

Note that all but the first of the BB solutions of Table 5 are dominated by at least one other solution which took up to the same time in providing a solution of at least the same quality. Fig. 4 is a scatter graph of the solutions that took up to 110 min and shows the efficient frontier (Goodwin & Wright, 1998)

Table 5

Summary of heuristic solutions and computing times (in decreasing order of solution value)

Heuristic	Mean solution	Mean time (min)
Section 3.3: backward-then-forward pass	505,000	0.03
Table 1: default BB search for 1 min	481,713	1
Table 1: default BB search for 20 min	430,532	20
Table 3: 100 iterations of descent algorithm for a^*	419,812	16.5
Table 1: default BB search for 30 min	414,358	30
Table 4: 1000 iterations of DN search on z_{lt} (av.)	414,000	16
Table 4: 1000 iterations of descent algorithm for z_{lt}	411,855	16
Table 4: 10,000 iterations of descent algorithm for z_{lt}	411,271	110
Table 3: 100 iterations of DN search on a^*	408,965	16.5
Table 1: default BB search for 90 min	400,408	90
Table 1: default BB search for 300 min	400,408	300
Table 3: 1000 iterations of DN search on a^*	404,549	169
Table 4: 10,000 iteration of DN search on z_{lt}	394,006	110

that links all the non-dominated solutions, providing a user with several alternatives depending on how much time is available to obtain a result. If an instant solution is needed, then use the forward-then-backward pass of Section 3.3. If just a few minutes are available, then default BB will suffice. The 120° lower left corner of the efficient frontier wraps a tight cluster of three nearly equal solutions obtainable after 16 or 17 min with DN search on a^* performing best. These are the ‘best value for money’ solutions in that only slightly better BB and DNS solutions are obtainable with much more computing time.

The a^* approach performs well with DN search. However, further research is needed to investigate its behaviour and performance on more complex problems, such as multiple non-identical machines where

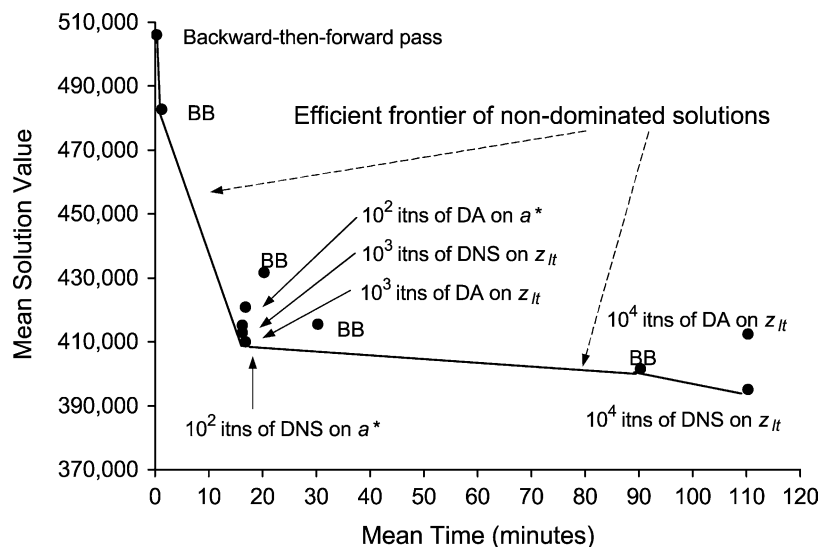


Fig. 4. Efficient frontier of the solutions and times of Table 5.

the unit production time a varies with product and machine. In this case, the local search becomes more challenging. DN search and the descent algorithm on the tank binary variables z_{lt} also perform well, with the descent algorithm tending to get stuck in a local optimum after 3000 iterations.

The research of this paper was carried out after the termination of the project at the drinks manufacturer, but the model and data are typical of many production planning challenges encountered in practice. The purpose of the paper submission was to inform both the academic and practitioner industrial engineering community of the interesting and useful results. These indicate that hybrid heuristics mixing both classical BB methods and modern local search approaches are worth exploring to accelerate a search for good solutions in industrial production planning and scheduling.

References

- Aarts, E. H. L., Korst, J. H. M., & van Laarhoven, P. J. M. (1997). Simulated annealing. In E. H. L. Aarts, & J. K. Lenstra (Eds.), *Local search in combinatorial optimization* (pp. 91–120). Chichester: Wiley.
- Aarts, E. H. L., & Lenstra, J. K. (Eds.), (1997). *Local search in combinatorial optimization*. Chichester: Wiley.
- Ahuja, R. K., Ergun, O. K., Orlin, J. B., & Punnen, A. P. (1999). A survey of very large neighbourhood search techniques. *Working paper*, Operations Research Center, Massachusetts Institute of Technology.
- Belvaux, G. & Wolsey, L. A. (1998). Lot-sizing problems: Modelling issues and a specialized branch-and-cut system *BC-PROD*, CORE Discussion Paper 9849, Université Catholique de Louvain, Belgium.
- Dowsland, K. A. (1993). Simulated annealing. In C. R. Reeves (Ed.), *Modern heuristic techniques for combinatorial problems* (pp. 20–69). London: Blackwell Scientific, chapter 2.
- Dueck, G., & Scheuer, T. (1990). Threshold accepting: A general purpose optimization algorithm appearing superior to simulated annealing. *Journal of Computational Physics*, 90(1), 161–175.
- Glass, C., & Potts, C. N. (1996). A comparison of local search methods for flow shop scheduling. *Annals of Operations Research*, 63, 489–509.
- Glover, F., & Laguna, M. (1997). *Tabu search*. Boston: Kluwer Academic Publishers.
- Goodwin, P., & Wright, G. (1998). *Decision analysis for management judgement* (2nd ed.). Chichester: Wiley.
- Hansen, P., & Mladenovic, N. (2001). Variable neighborhood search: Principles and applications. *European Journal of Operational Research*, 130, 449–467.
- Hertz, A., Taillard, E., & de Werra, D. (1997). Tabu search. In E. H. L. Aarts, & J. K. Lenstra (Eds.), *Local search in combinatorial optimization* (pp. 121–136). Chichester: Wiley.
- Lin, C. K. Y., Haley, K. B., & Sparks, C. (1995). A comparative study of both standard and adaptive versions of threshold accepting and simulated annealing algorithms in 3 scheduling problems. *European Journal of Operational Research*, 83(2), 330–346.
- Meyr, H. (2000). Simultaneous lot-sizing and scheduling by combining local search with dual optimization. *European Journal of Operational Research*, 120, 311–326.
- Mühlenbein, H., & Zimmerman, J. (2000). Size of neighbourhood more important than temperature for stochastic local search. *Proceedings of the 2000 Congress on Evolutionary Computation, July 16–19, 2000, La Jolla, CA* (pp. 1017–1024). Piscataway, NJ: IEEE.
- Pinedo, M., & Chao, X. (1999). *Operations scheduling with applications in manufacturing and services*. Irwin/McGraw-Hill.
- Pirlot, M. (1996). General local search methods. *European Journal of Operational Research*, 92, 493–511.
- Reeves, C. (1994). Genetic algorithms and neighbourhood search. In T. C. Fogarty (Ed.), *Evolutionary computing: ASIB workshop, Leeds, UK: Selected papers*. Berlin: Springer.
- Sait, S. M., & Youssef, H. (1999). *Iterative computer algorithms with applications in engineering*. Los Alamitos, CA: IEEE Computer Society.
- Wolsey, L. A. (1998). *Integer programming*. New York: Wiley.